

R Bookclub

Chapters 2 & 3

Shu Asai

04/03/2019

Chapter 2: Workflow basics

R as a calculator

```
2+4
```

```
[1] 6
```

```
pi
```

```
[1] 3.141593
```

```
(pi+1)/2
```

```
[1] 2.070796
```

Objects & Naming

Use '<-' to attribute value to objects

```
floors_harwick <- 8
```

```
floors_gonda <- 19
```

Exercise:

```
floor_diff <- floors_gonda - floors_harwick
```

```
floor_diff
```

```
[1] 11
```

Functions

Provide input arguments, Obtain output

Example: **seq()**

- ▶ Two arguments needed: `seq(from, to)`
- ▶ R produces a **sequence** of integers from provided arguments

```
seq(10,20)
```

```
[1] 10 11 12 13 14 15 16 17 18 19 20
```

Function documentation is found by typing “?” + function name in the Console:


```
?seq
```

Exercises

Find the error in my code

```
ggplot(dota = mpg) +  
  geom_point(mapping = aes(x=displ, y=hwy))
```

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x=displ, y=hwy))
```



Exercises cont.

```
fliter(mpg, cyl=8)
```

```
filter(mpg, cyl==8)
```

```
# A tibble: 70 x 11
```

	manufacturer	model	displ	year	cyl	trans	drv	cty	hwy	fl	class
	<chr>	<chr>	<dbl>	<int>	<int>	<chr>	<chr>	<int>	<int>	<chr>	<chr>
1	audi	a6 q~	4.2	2008	8	auto~	4	16	23	p	mids
2	chevrolet	c150~	5.3	2008	8	auto~	r	14	20	r	suv
3	chevrolet	c150~	5.3	2008	8	auto~	r	11	15	e	suv
4	chevrolet	c150~	5.3	2008	8	auto~	r	14	20	r	suv
5	chevrolet	c150~	5.7	1999	8	auto~	r	13	17	r	suv
6	chevrolet	c150~	6	2008	8	auto~	r	12	17	r	suv
7	chevrolet	corv~	5.7	1999	8	manu~	r	16	26	p	2sea
8	chevrolet	corv~	5.7	1999	8	auto~	r	15	23	p	2sea
9	chevrolet	corv~	6.2	2008	8	manu~	r	16	26	p	2sea
10	chevrolet	corv~	6.2	2008	8	auto~	r	15	25	p	2sea

```
# with 60 more rows
```

Exercises cont.

```
filter(diamond, carat>3)
```

```
filter(diamonds, carat>3)
```

```
# A tibble: 32 x 10
```

	carat	cut	color	clarity	depth	table	price	x	y	z
	<dbl>	<ord>	<ord>	<ord>	<dbl>	<dbl>	<int>	<dbl>	<dbl>	<dbl>
1	3.01	Premium	I	I1	62.7	58	8040	9.1	8.97	5.67
2	3.11	Fair	J	I1	65.9	57	9823	9.15	9.02	5.98
3	3.01	Premium	F	I1	62.2	56	9925	9.24	9.13	5.73
4	3.05	Premium	E	I1	60.9	58	10453	9.26	9.25	5.66
5	3.02	Fair	I	I1	65.2	56	10577	9.11	9.02	5.91
6	3.01	Fair	H	I1	56.1	62	10761	9.54	9.38	5.31
7	3.65	Fair	H	I1	67.1	53	11668	9.53	9.48	6.38
8	3.24	Premium	H	I1	62.1	58	12300	9.44	9.4	5.85
9	3.22	Ideal	I	I1	62.6	55	12545	9.49	9.42	5.92
10	3.5	Ideal	H	I1	62.8	57	12587	9.65	9.59	6.03

```
# with 22 more rows
```


Keyboard Tips & Shortcuts

- ▶ Snake Case (Examples: `"floors_harwick"` , `"cohort_A"` , `"scatter_MCHS"`)
 - * Descriptive and readable object names
 - * Avoid camel case and periods

Keyboard Tips & Shortcuts

- ▶ Snake Case (Examples: “floors_harwick” , “cohort_A” , “scatter_MCHS”)
 - * Descriptive and readable object names
 - * Avoid camel case and periods
- ▶ Use **(Alt)** **(-)** as a keyboard shortcut for `<-`

Keyboard Tips & Shortcuts

- ▶ Snake Case (Examples: `"floors_harwick"` , `"cohort_A"` , `"scatter_MCHS"`)
 - * Descriptive and readable object names
 - * Avoid camel case and periods
- ▶ Use **(Alt) (-)** as a keyboard shortcut for `<-`
- ▶ Using "Tab" to explore objects

Keyboard Tips & Shortcuts

- ▶ Snake Case (Examples: “floors_harwick” , “cohort_A” , “scatter_MCHS”)
 - * Descriptive and readable object names
 - * Avoid camel case and periods
- ▶ Use **(Alt) (-)** as a keyboard shortcut for <-
- ▶ Using “Tab” to explore objects
- ▶ RStudio Environment & History: (Cmd) (Ctrl) (-) (up arrow)

Keyboard Tips & Shortcuts

- ▶ Snake Case (Examples: “floors_harwick” , “cohort_A” , “scatter_MCHS”)
 - * Descriptive and readable object names
 - * Avoid camel case and periods
- ▶ Use **(Alt) (-)** as a keyboard shortcut for <-
- ▶ Using “Tab” to explore objects
- ▶ RStudio Environment & History: (Cmd) (Ctrl) (-) (up arrow)
- ▶ Shortcut Cheatsheet: (Alt) (Shift) (K)

Chapter 3: Transformations with dplyr

Variable Types

- ▶ *int* (Integers)
- ▶ *dbl* (Doubles or real numbers)
- ▶ *chr* (Character vectors or strings)
- ▶ *dtm* (Date / Times)

?dplyr

What is it?

dplyr is an R package containing functions useful to transform and manipulate data.

Why dplyr?

Intuitive, readable, and fast!

Five key dplyr functions:

1. filter
2. arrange
3. select
4. mutate
5. summarize

dplyr Function Workflow

Load in the dplyr package

```
library(dplyr)
```

1. Provide the function a dataframe

dplyr Function Workflow

Load in the dplyr package

```
library(dplyr)
```

1. Provide the function a dataframe
2. Describe what to do to the dataframe with column specifications

dplyr Function Workflow

Load in the dplyr package

```
library(dplyr)
```

1. Provide the function a dataframe
2. Describe what to do to the dataframe with column specifications
3. Output to a new dataframe

The Data

flights

A tibble: 336,776 x 19

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838
10	2013	1	1	558	600	-2	753

... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,

arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,

'filter': Subset Rows

'filter' workflow

- ▶ Provide the function a dataframe

```
filter(flights)
```

- ▶ Describe what to do to the dataframe with column specifications

```
filter(flights, month == 1 , day == 1)
```

- ▶ Output to a new dataframe

```
(flights_1_1 <- filter(flights, month == 1 , day == 1))
```

```
# A tibble: 842 x 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004

Useful Operators

Use operators to specify your “filter” statement.

1. “and” statement: **&**

Find all flights operated by United (“UA”), American (“AA”), or Delta (“DL”)

```
filter(flights, carrier %in% c( "UA", "AA", "DL" ))
```

Useful Operators

Use operators to specify your “filter” statement.

1. “and” statement: `&`
2. “or” statement: `|`

Find all flights operated by United (“UA”), American (“AA”), or Delta (“DL”)

```
filter(flights, carrier %in% c( "UA", "AA", "DL" ))
```


Useful Operators

Use operators to specify your “filter” statement.

1. “and” statement: `&`
2. “or” statement: `|`
3. “not” statement: `!=`

Find all flights operated by United (“UA”), American (“AA”), or Delta (“DL”)

```
filter(flights, carrier %in% c( "UA", "AA", "DL" ))
```

Useful Operators

Use operators to specify your “filter” statement.

1. “and” statement: **&**
2. “or” statement: **|**
3. “not” statement: **!=**
4. extended “or” statement: **%in%**

Find all flights operated by United (“UA”), American (“AA”), or Delta (“DL”)

```
filter(flights, carrier %in% c( "UA", "AA", "DL" ))
```

More Exercises

Find flights that had an arrival delay of two or more hours:

```
filter(flights, arr_delay >=2)
```

Find flights that flew to "IAH" or "HOU"

```
filter(flights, dest == "IAH" | dest == "HOU")  
filter(flights, dest %in% c("IAH", "HOU"))
```

Find flights that were delayed by at least an hour, but made up over 30 minutes.

```
filter(flights, dep_delay >= 1 & arr_delay < -30)
```

Exercises cont.

The *between(column, min, max)* function is also a useful dplyr verb.

```
filter( dataframe, between(column, min, max) )
```

How can we use it in this example?

Find flights that departed between July, August, and September.

```
filter(flights, between(month,7,9))
```

```
# A tibble: 3 x 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	7	1	1	2029	212	236
2	2013	7	1	2	2359	3	344
3	2013	7	1	29	2245	104	151

```
# ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,  
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
#   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
```

Exercises cont.

Which flights are missing a "dep_time"?

```
filter(flights, is.na(dep_time))
```

```
# A tibble: 3 x 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	NA	1630	NA	NA
2	2013	1	1	NA	1935	NA	NA
3	2013	1	1	NA	1500	NA	NA

```
# ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,  
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
#   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
#   time_hour <dtm>
```

What could these rows represent?

'arrange': Order your Data

'arrange'

Order dataset by provided columns in ascending order

```
(flight_month <- arrange(flights, month))
```

A tibble: 336,776 x 19

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838
10	2013	1	1	558	600	-2	753

... with 336,766 more rows, and 12 more variables: sched_arr_time <int>

'arrange' cont.

Arranging the data in the opposite direction (descending order):

```
(flight_year_month <- arrange(flights, desc(month)))
```

A tibble: 336,776 x 19

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	12	1	13	2359	14	446
2	2013	12	1	17	2359	18	443
3	2013	12	1	453	500	-7	636
4	2013	12	1	520	515	5	749
5	2013	12	1	536	540	-4	845
6	2013	12	1	540	550	-10	1005
7	2013	12	1	541	545	-4	734
8	2013	12	1	546	545	1	826
9	2013	12	1	549	600	-11	648
10	2013	12	1	550	600	-10	825

... with 336,766 more rows, and 12 more variables: sched_arr_time <int>

'arrange': cont.

Multiple arguments in the arrange statement.

```
flight_year_dep <- arrange(flights, year, dep_time)
```

```
# A tibble: 4 x 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	13	1	2249	72	108
2	2013	1	31	1	2100	181	124
3	2013	11	13	1	2359	2	442
4	2013	12	16	1	2359	2	447

```
# ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,  
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
#   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
#   time_hour <dtm>
```

What happens to NA values in arrange() statements?

Handling NAs and Missing Values

To determine if a value is missing/NA, use the 'is.na()' function:

```
x <- NA  
is.na(x)
```

```
[1] TRUE
```

The following result in NA:

```
NA > 5  
10 == NA  
NA + 10  
NA / 2
```

Rule- if the calculation involves NA, result will almost always be NA.

NAs: Exceptions

```
NA ^ 0
```

```
[1] 1
```

```
NA | TRUE
```

```
[1] TRUE
```

Exercises

How can we sort missing values to the top of the dataset?

```
arrange(flights, desc(is.na(dep_time)))
```

Find flights that were most delayed, then sort by those that left earliest

```
arrange(flights, desc(dep_delay), departure)
```

Find the longest flights

```
arrange(flights, desc(air_time))
```

'select': Select Columns of Data

'select'

```
flights_ymd <- select(flights, year, month, day)
flights_ymd
```

```
# A tibble: 336,776 x 3
```

	year	month	day
	<int>	<int>	<int>
1	2013	1	1
2	2013	1	1
3	2013	1	1
4	2013	1	1
5	2013	1	1
6	2013	1	1
7	2013	1	1
8	2013	1	1
9	2013	1	1
10	2013	1	1

```
# ... with 336,766 more rows
```

'select' cont.

Selecting all columns between 'carrier' and 'origin'

```
select(flights, carrier:origin)
```

```
# A tibble: 336,776 x 4
```

	carrier	flight	tailnum	origin
	<chr>	<int>	<chr>	<chr>
1	UA	1545	N14228	EWR
2	UA	1714	N24211	LGA
3	AA	1141	N619AA	JFK
4	B6	725	N804JB	JFK
5	DL	461	N668DN	LGA
6	UA	1696	N39463	EWR
7	B6	507	N516JB	EWR
8	EV	5708	N829AS	LGA
9	B6	79	N593JB	JFK
10	AA	301	N3ALAA	LGA

```
# ... with 336,766 more rows
```

'select' cont.

Selecting all columns except those from 'carrier' to 'origin'

```
select(flights, -(carrier : origin))
```

A tibble: 336,776 x 15

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838
10	2013	1	1	558	600	-2	753

... with 336,766 more rows, and 8 more variables: sched_arr_time <int>.

'select' helper statements

starts_with("x")

selects columns that begin with "x"

ends_with("type")

selects columns that end with "type"

contains("gene")

selects columns that contain "gene"

num_range("x", 1:3)

selects columns titled, "x1", "x2", and "x3"

Exercises

What happens if you refer to the same column multiple times in a 'select' call?

```
select(flights, year, year)
```

```
# A tibble: 3 x 1
```

```
  year
```

```
<int>
```

```
1  2013
```

```
2  2013
```

```
3  2013
```

Exercises cont.

What does the 'one_of' function do? How could it be useful in a select statement?

Given a vector of characters, 'one_of' finds column names that match in the vector.

```
variables <- c("year", "month", "day", "arr_time")
```

```
select(flights, one_of(variables))
```

```
# A tibble: 3 x 4
```

	year	month	day	arr_time
	<int>	<int>	<int>	<int>
1	2013	1	1	830
2	2013	1	1	850
3	2013	1	1	923

Column Manipulation

Use 'rename' to change column names

Syntax: `rename(data, new_column_name = old_column_name)`

```
rename(flights, depart_delay = dep_delay)
```

A tibble: 336,776 x 19

	year	month	day	dep_time	sched_dep_time	depart_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838

Column Manipulation cont.

Use 'select' and 'everything' together to rearrange columns

```
select(flights, time_hour, air_time, everything())
```

```
# A tibble: 336,776 x 19
```

	time_hour	air_time	year	month	day	dep_time	sched_dep_time
	<dtm>	<dbl>	<int>	<int>	<int>	<int>	<int>
1	2013-01-01 05:00:00	227	2013	1	1	517	515
2	2013-01-01 05:00:00	227	2013	1	1	533	529
3	2013-01-01 05:00:00	160	2013	1	1	542	540
4	2013-01-01 05:00:00	183	2013	1	1	544	545
5	2013-01-01 06:00:00	116	2013	1	1	554	600
6	2013-01-01 05:00:00	150	2013	1	1	554	558
7	2013-01-01 06:00:00	158	2013	1	1	555	600
8	2013-01-01 06:00:00	53	2013	1	1	557	600
9	2013-01-01 06:00:00	140	2013	1	1	557	600
10	2013-01-01 06:00:00	138	2013	1	1	558	600

```
# ... with 336,766 more rows, and 12 more variables: dep_delay <dbl>.
```

mutate: Create Columns

'mutate'

Create and append columns to existing data with 'mutate'

Syntax:

```
mutate(data, new_column_name = column manipulation )
```

'mutate' example

```
flights_A <- select(flights, year : day, arr_delay, starts_with("dep"))
```

```
# A tibble: 3 x 6
```

	year	month	day	arr_delay	dep_time	dep_delay
	<int>	<int>	<int>	<dbl>	<int>	<dbl>
1	2013	1	1	11	517	2
2	2013	1	1	20	533	4
3	2013	1	1	33	542	2

```
mutate(flights_A,  
       gain = arr_delay - dep_delay)
```

```
# A tibble: 3 x 7
```

	year	month	day	arr_delay	dep_time	dep_delay	gain
	<int>	<int>	<int>	<dbl>	<int>	<dbl>	<dbl>
1	2013	1	1	11	517	2	9
2	2013	1	1	20	533	4	16
3	2013	1	1	33	542	2	31

'mutate' Multiple Columns

```
mutate(flights_A,  
       gain = arr_delay - dep_delay,  
       arr_dep_diff = arr_delay - dep_time)
```

```
# A tibble: 5 x 8
```

	year	month	day	arr_delay	dep_time	dep_delay	gain	arr_dep_diff
	<int>	<int>	<int>	<dbl>	<int>	<dbl>	<dbl>	<dbl>
1	2013	1	1	11	517	2	9	-506
2	2013	1	1	20	533	4	16	-513
3	2013	1	1	33	542	2	31	-509
4	2013	1	1	-18	544	-1	-17	-562
5	2013	1	1	-25	554	-6	-19	-579

'transmute'

'transmute' is the same as 'mutate', but will instead only keep the new variables:

```
transmute(flights_A,  
  gain = arr_delay - dep_delay,  
  arr_dep_diff = arr_delay - dep_time)
```

A tibble: 336,776 x 2

	gain	arr_dep_diff
	<dbl>	<dbl>
1	9	-506
2	16	-513
3	31	-509
4	-17	-562
5	-19	-579
6	16	-542
7	24	-536
8	-11	-571
9	-5	-565

Useful Functions When Mutating

$/$, $+$, $-$, $*$, $^$

sum() , *mean()*

%/% (integer division)

```
4 %/% 2
```

```
[1] 2
```

%% (remainder)

```
4 %% 2
```

```
[1] 0
```

Useful Functions When Mutating cont.

Logs

- ▶ Natural logarithms

```
log()
```

- ▶ Binary logarithms (base 2)

```
log2()
```

- ▶ Common logarithms (base 10)

```
log10()
```

Useful Functions cont : Offsets

'lead' & 'lag'

```
(tn <- 1:10)
```

```
[1]  1  2  3  4  5  6  7  8  9 10
```

```
(lead(tn))
```

```
[1]  2  3  4  5  6  7  8  9 10 NA
```

```
(lag(tn))
```

```
[1] NA  1  2  3  4  5  6  7  8  9
```

Useful Functions cont : Aggregates

Cumulative and rolling calculations

What is the sum of all previous values (including current value)?

```
cumsum(tn)
```

```
[1]  1  3  6 10 15 21 28 36 45 55
```

What is the mean of all previous values (including current value)?

```
cummean(tn)
```

```
[1] 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0 5.5
```

others: 'cummin' , 'cummax', 'cumprod'

Useful Functions cont : Comparisons

Logical comparisons

<

>

<=

>=

!=

==

Exercises

Find the 10 most delayed flights using mutate and the 'min_rank' function.

Note: 'min_rank' assigns a number equal to a number of elements less than that value plus one.

```
mr <- c(4, 2, 7, 7, 7, 0, 10)  
min_rank(mr)
```

```
[1] 3 2 4 4 4 1 7
```

```
min_rank(desc(mr))
```

```
[1] 5 6 2 2 2 7 1
```


Exercises cont.

Find the 10 most delayed flights using mutate and the 'min_rank' function.

```
flights_delay <- mutate(flights, delay_rank = min_rank(desc(dep_delay)))
```

```
flights_delay <- select(flights_delay, delay_rank, dep_delay, everything())
```

```
arrange(flights_delay, delay_rank)[1:5,]
```

```
# A tibble: 5 x 20
```

	delay_rank	dep_delay	year	month	day	dep_time	sched_dep_time	arr_time
	<int>	<dbl>	<int>	<int>	<int>	<int>	<int>	<int>
1	1	1301	2013	1	9	641	900	1242
2	2	1137	2013	6	15	1432	1935	1607
3	3	1126	2013	1	10	1121	1635	1239
4	4	1014	2013	9	20	1139	1845	1457
5	5	1005	2013	7	22	845	1600	1044

```
# ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
```

```
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
```

Exercises cont.

'summarize': Aggregate Calculations

'summarize'

Collapses data into a single row:

```
summarize(flights, avg_delay = mean(dep_delay, na.rm = TRUE))
```

```
# A tibble: 1 x 1
```

```
  avg_delay
```

```
    <dbl>
```

```
1      12.6
```

'summarize' is not too useful on its own.

It is often used in tandem with

Pipes and 'group_by'

Pipes and 'group_by'

Grouped Summaries with “group_by”

Implement *group_by* to use dplyr functions on a *grouped* dataframe

Example:

```
by_day <- group_by(flights, year, month, day)
```

The above code does nothing to the structure of the data. The data is merely coerced into a grouped dataframe.

```
summarize(by_day, delay = mean(dep_delay, na.rm = TRUE))[1:3,]
```

```
# A tibble: 3 x 4
```

```
# Groups:   year, month [1]
```

	year	month	day	delay
	<int>	<int>	<int>	<dbl>
1	2013	1	1	11.5
2	2013	1	2	13.9
3	2013	1	3	11.0

A note on 'na.rm'

Recall, *most* calculations that include 'NA' values will result in an 'NA'.

```
mean(c(2,3,4,NA, 6))
```

```
[1] NA
```

Omit the 'NA' values in the calculation with use the option *na.rm=TRUE*

```
mean(c(2,3,4,NA,6), na.rm = TRUE)
```

```
[1] 3.75
```

Ungrouping

Data can be ungrouped at any time

```
by_day <- group_by(flights, year, month, day)
```

```
orig_data <- ungroup(by_day)
```


Piping with '%>%'

Pipes allow multiple 'dplyr' functions to be used consecutively.

For example, you may want to take your original data, and then:

1. Select specific columns
2. Filter for observations
3. Create (mutate) another column

Data %>% Select columns %>% Filter observations %>% Create a column

Pipe Framework

```
flight_a <- flights %>% select(year, air_time, dest) %>% mutate(air_time_h
```

1. Provide a dataset

```
# A tibble: 3 x 4
  year air_time dest  air_time_hours
  <int>   <dbl> <chr>         <dbl>
1  2013     227 IAH             3.78
2  2013     227 IAH             3.78
3  2013     160 MIA             2.67
```

Pipe Framework

```
flight_a <- flights %>% select(year, air_time, dest) %>% mutate(air_time_h
```

1. Provide a dataset
2. Pipe (%>%)

```
# A tibble: 3 x 4
```

	year	air_time	dest	air_time_hours
	<int>	<dbl>	<chr>	<dbl>
1	2013	227	IAH	3.78
2	2013	227	IAH	3.78
3	2013	160	MIA	2.67

Pipe Framework

```
flight_a <- flights %>% select(year, air_time, dest) %>% mutate(air_time_h
```

1. Provide a dataset
2. Pipe (%>%)
3. Input desired 'dplyr' function

```
# A tibble: 3 x 4
```

	year	air_time	dest	air_time_hours
	<int>	<dbl>	<chr>	<dbl>
1	2013	227	IAH	3.78
2	2013	227	IAH	3.78
3	2013	160	MIA	2.67

Pipe Framework

```
flight_a <- flights %>% select(year, air_time, dest) %>% mutate(air_time_h
```

1. Provide a dataset
2. Pipe (%>%)
3. Input desired 'dplyr' function
4. Repeat steps 2 - 3 as needed

```
# A tibble: 3 x 4
```

	year	air_time	dest	air_time_hours
	<int>	<dbl>	<chr>	<dbl>
1	2013	227	IAH	3.78
2	2013	227	IAH	3.78
3	2013	160	MIA	2.67

Pipe Framework

```
flight_a <- flights %>% select(year, air_time, dest) %>% mutate(air_time_h
```

1. Provide a dataset
2. Pipe (%>%)
3. Input desired 'dplyr' function
4. Repeat steps 2 - 3 as needed
5. Output into a new object

```
# A tibble: 3 x 4
```

	year	air_time	dest	air_time_hours
	<int>	<dbl>	<chr>	<dbl>
1	2013	227	IAH	3.78
2	2013	227	IAH	3.78
3	2013	160	MIA	2.67

Pipes and 'group_by' together

Using both pipes and 'group_by' allows data manipulation to groups of data.

Consider the question: What was the average departure delay in each month?

'group_by' month, use the 'summarize' function

```
flights %>% group_by(month) %>% summarize(avg_delay = mean(dep_delay, na.rm=
```

```
# A tibble: 4 x 2
```

```
  month avg_delay
```

```
  <int>   <dbl>
```

```
1     1    10.0
```

```
2     2    10.8
```

```
3     3    13.2
```

```
4     4    13.9
```

Useful Summary Functions

```
flights %>% group_by(month) %>% summarize()
```

Location: *mean(x), median(x)*

Spread: *sd(x), IQR(x), mad(x) (median absolute deviation)*

Rank: *min(x), quantile(x, 0.5) , max(x)*

Position: *first(x) , nth(x, 5), last(x)*

Counts: *n(), sum(is.na(x)), sum(!(is.na(x))), n_distinct(x)*

Find the sum of 'NAs' in the origin column in each month:

```
flights %>% group_by(month) %>% summarize(na_origin = sum(is.na(origin)))
```


Exercises

Provide another approach to attain output from the following:

```
not_cancelled %>% count(dest)
```

Essentially: Find all flights that are not cancelled.

One solution:

```
not_cancelled <- flights %>%  
  filter(!is.na(dep_delay), !is.na(arr_delay))
```

Other solutions?

If a flight never departs, then it won't arrive. But a flight could also depart and not arrive (crashes, lost flights). We could use 'arr_delay' as a proxy to define cancelled flights.

Exercises cont.

For each plane, count the number of flights before the first delay of greater than one hour.

What dplyr verbs do we want to use?

```
flights %>%  
  arrange(tailnum, year, month, day) %>%  
  group_by(tailnum) %>%  
  mutate(delay_gt1hr = dep_delay > 60) %>%  
  mutate(before_delay = cumsum(delay_gt1hr)) %>%  
  filter(before_delay < 1) %>%  
  count(sort = TRUE)
```

Exercises

For each destination, compute the total minutes of delay.

For each flight, compute the proportion of the total delay for its destination.

```
flights %>%  
  filter(!is.na(arr_delay), arr_delay > 0) %>%  
  group_by(dest) %>%  
  mutate(  
    arr_delay_total = sum(arr_delay),  
    arr_delay_prop = arr_delay / arr_delay_total) %>%  
  ungroup()
```

Exercises cont.

Find destinations that are flown by at least two carriers.

```
dest_2carriers <- flights %>%  
  # keep only unique carrier,dest pairs  
  select(dest, carrier) %>%  
  group_by(dest, carrier) %>%  
  filter(row_number() == 1) %>%  
  # count carriers by destination  
  group_by(dest) %>%  
  mutate(n_carrier = n_distinct(carrier)) %>%  
  filter(n_carrier >= 2)
```